



HaaS with Remote Partial Configuration

A Remote Partially Reconfigurable RISC-V
Soft Processor on a Low end FPGA

Table of Content

- Introduction
- Technical
- Similar Work
- The Proposal
- Verification
- Conclusion
- Links



Introduction

- Imagine a data center with a massive stack of FPGAs providing HaaS
- To efficiently use FPGAs, we need virtually no downtime
- Partial reconfiguration provides an optimal solution to such a venture
- There are some roadblocks to utilise such an architecture efficiently



Contd.

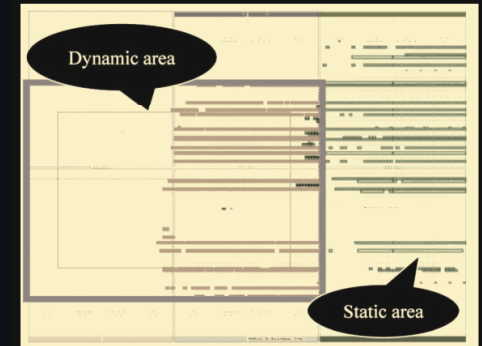
- To implement partial configuration on an FPGA, the usual approach is to pair the FPGA fabric with a processor (like ARM cortex / RISC-V core) as a controller
- This increases the cost of individual Units \therefore affecting scalability
- Alternatively we can have the FPGA fabric separated and controlled via JTAG cables and the server acting as the controller
- But this is unfeasible because the partial bitstreams are sensitive to noise which is exacerbated along long distances
- The solution proposed in this paper is to divide the FPGA fabric into dynamic and static regions and implement the PR (Partial Reconfiguration) Controller in the static partition



Technical

Partial Bitstreams

- Partial bitstreams are the bitstreams created by the EDA tool to implement Partial Configuration on an FPGA
- After compatibility checks with the static modules. This file is then configured onto the FPGA fabric using ICAP (Internal Configuration Access Point)



LiteX

- Generates a source code for a riscv soft core
- This SoC can be implemented in any FPGA

Similar Works

REoN

A network partial configuration framework that was apt for high end FPGAs. (Not suitable for low end FPGAs)

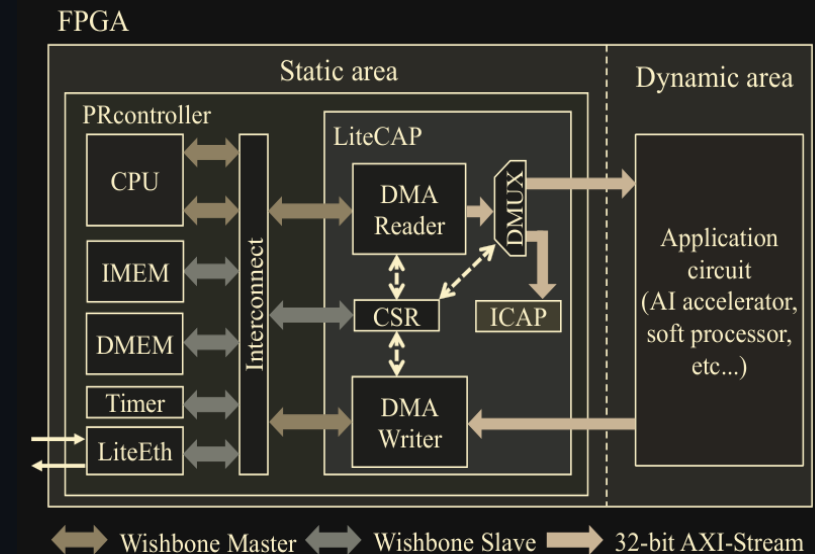
Hardware implementation of TCP/IP stack

But this module takes significant FPGA resources, which are limited in low end FPGAs



The Proposal

- Partition the FPGA fabric into static and dynamic sections
 - Static partition will contain the PR controller
 - Dynamic partition will contain the application circuit
- 2 TCP connections during operation between host and FPGA
 - One for sending the partial bitstream
 - Another to send the data to the application circuit after configuration



Verification

Environment

- Board used -> Arty A7-35T FPGA
- riscv64-unknown-elf-gcc 12.2.0
- LiteX
- VexRiscv -> RV32I

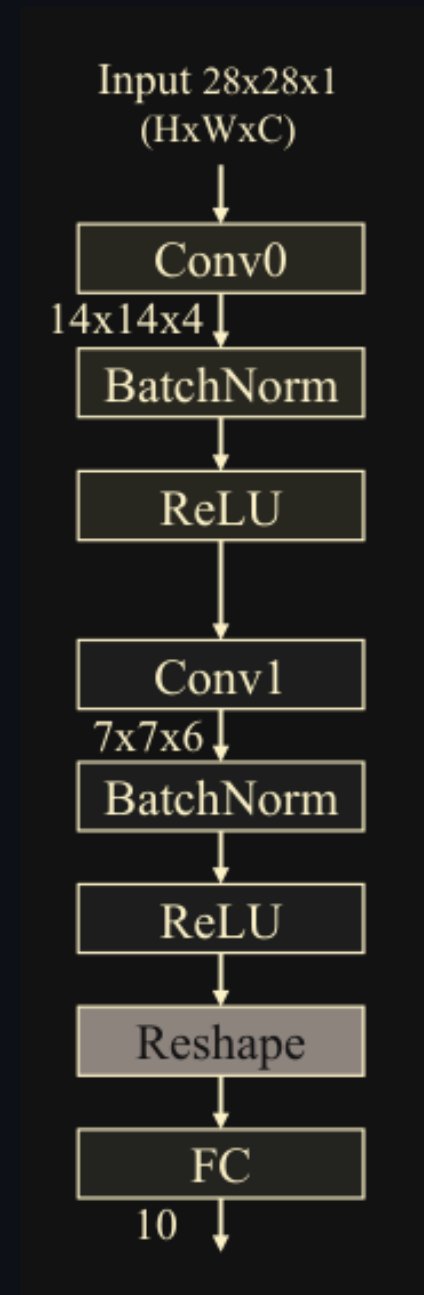
Tests

- CNN Network as Application Circuit
- RISCv Soft Processor as Application Circuit



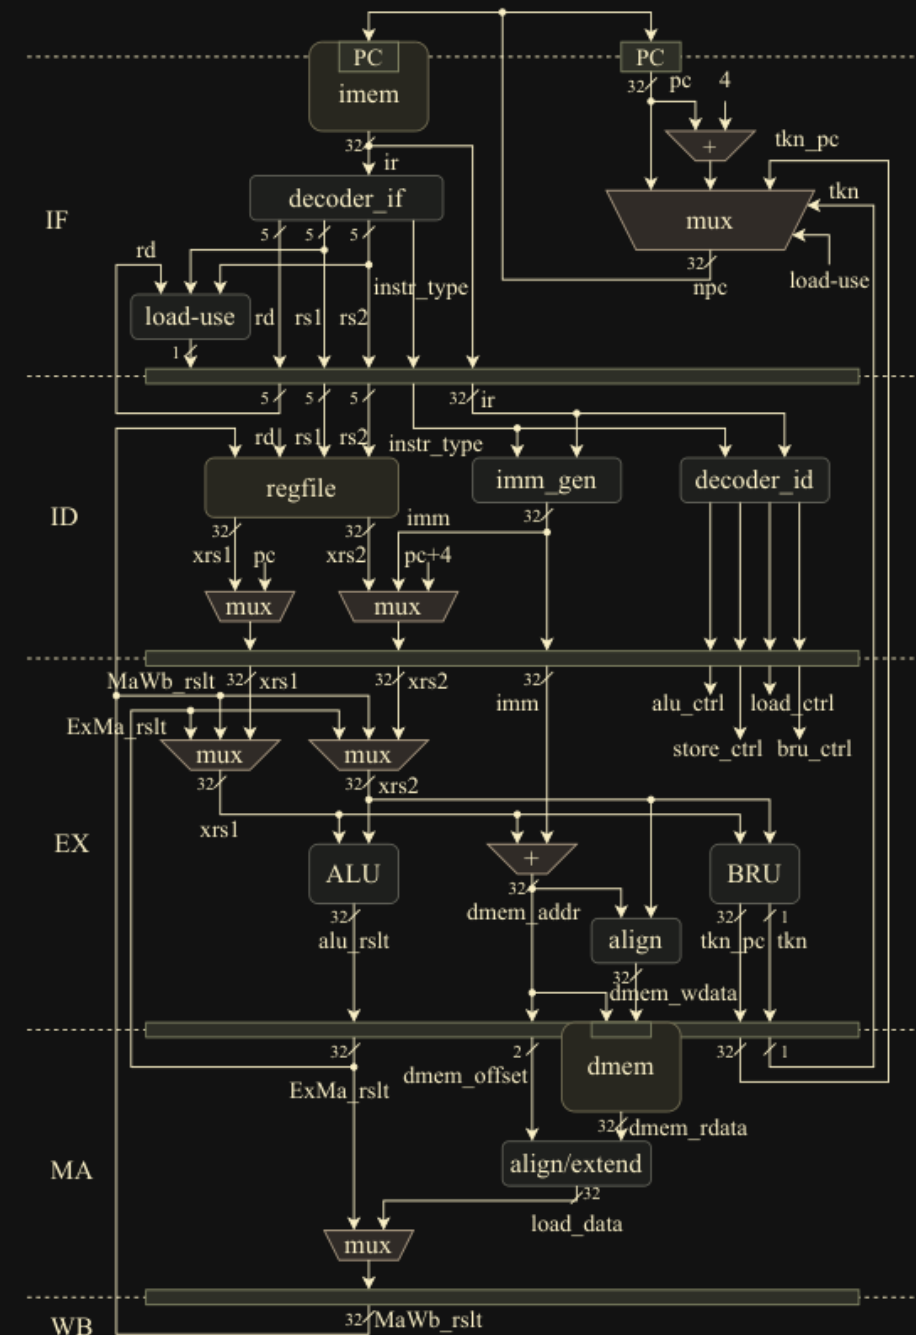
CNN Network as Application Circuit

- A handwriting classification model
- Database \rightarrow MNIST
- NN inference framework used to implement the CNN on FPGA \rightarrow FINN



RISCV Soft Processor as Application Circuit

- In-order execution model written in Verilog
- A Hello world program with output given to AXI-Stream interface



Verification Results


- To verify the reliability of the proposed SoC in real world applications, they connected the SoC and host PC to the laboratory LAN
- While sending the partial bitstreams continuously, 100 different MNIST images were send for inferencing. The result matched the HDL simulations
- Similarly for the soft core, while sending the partial bitstreams continuously, and sending the instructions for the "Hello, World!" program. There was no error under rigorous testing




Evaluation Results

Resource Utilisation


- CPU - With a CPU on board the FPGA
- REoN - Network protocol for partial configuration on high end FPGAs
- PR Controller - The static partition of our SoC



(a) LUT Utilization.



(b) FF Utilization.



(c) BRAM Utilization.

Fig. 7: The amount of hardware used by CPU of the proposed SoC, the whole static area, and REoN.

Evaluation Results

Throughput for different configurations

- 2MB CNN accelerator partial bistream including headers
- 2MB CNN accelerator partial bistream starting with sync word
- 10MB CNN accelerator partial bistream with 8 MB zero padding

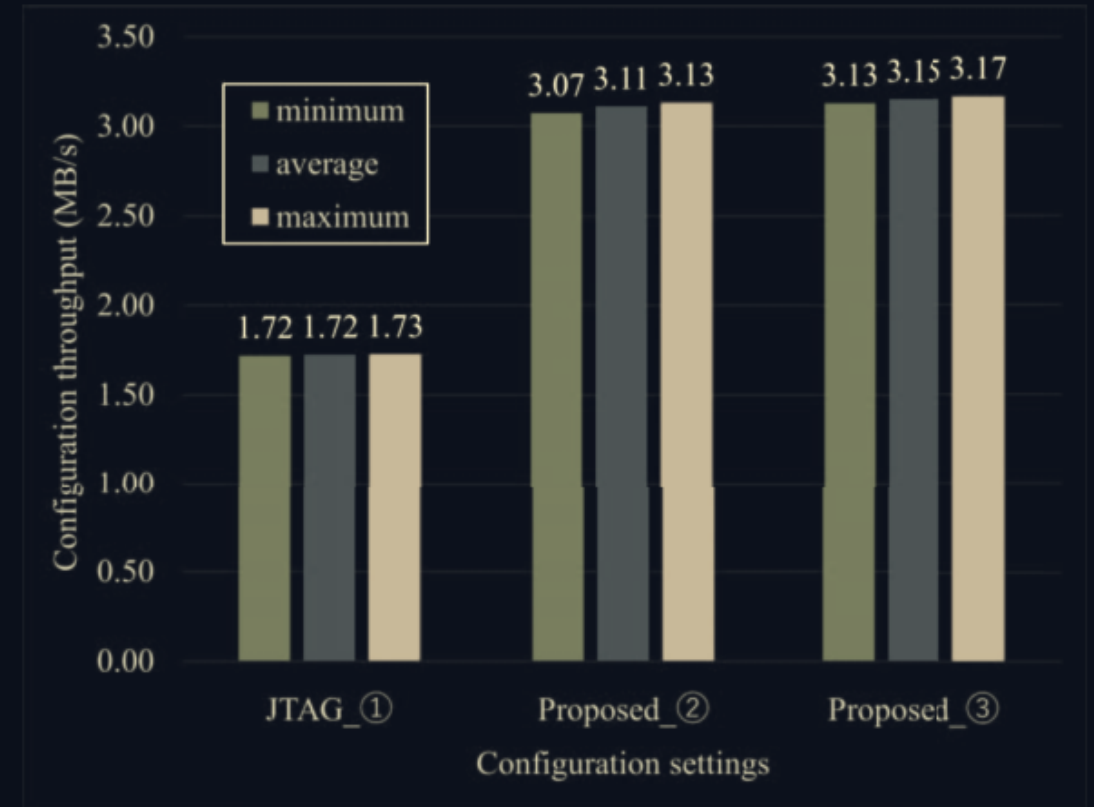


Fig. 8: The partial reconfiguration throughput in three settings.

Conclusion

- We use a stripped down version of other protocols and modules to fit it into a low end fpga with limited resources (Zynq vs Arty FPGA boards)
- The above proposal is suitable for datacenters and NFV (Network Function Virtualisation)
- The above SoC removes the need for an external CPU/ Controller that is usually used with partial reconfiguration of FPGAs. Making each FPGA standalone. Therefore more scalable



Links

Paper

- [The Main Paper](#)

Related Papers

- [REoN Protocol](#)
- [Limago: An FPGA-Based Open-Source 100 GbE TCP/IP Stack](#)
- [Scalable TCP/IP Stack Architecture for Reconfigurable Hardware](#)
- [IwIP](#)
- [RV CAP](#)



Thank You

